

UIPATH AGENTHACK 2026

# gemini-rpa-agent

An RPA failure-diagnosis agent that runs inside UiPath Maestro.

*Ask what broke. Get the failing step, the verbatim error, and the fix, in seconds.*

Maestro BPMN track

Gemini 2.5 + ADK

UiPath LLM Gateway

Apache-2.0

## THE PROBLEM

# The diagnosis is what burns the night, not the fix



**The fix is almost always copy-pasteable.** Finding it is what costs the night. `gemini-rpa-agent` walks the same tools an on-call walks, but in seconds, with the canonical fix already in hand.

## WHAT IT DOES

# One question in, a five-section diagnosis out

### YOU ASK

*“Workflow wf-onboarding-pro-2026-001 failed at 09:14 UTC. What broke, and how do I fix it?”*



### ANSWER

which workflow, which step, what broke

### EVIDENCE

verbatim error payload + step IDs, no paraphrasing

### ROOT CAUSE

one-sentence diagnosis

### REMEDIATION


copy-paste fix: old value, new value, retry command

### NEXT STEP

one follow-up check

## ARCHITECTURE

# The agent runs through the UiPath Platform, end to end



**Why this satisfies the rule.** Orchestration and agent logic run through the UiPath Platform: Maestro invokes the published agent, and every Gemini call passes through the UiPath LLM Gateway, not direct Vertex.

# Four read-only RPA tools, walked end to end

1

## `list_workflows(active_only)`

Workflows plus last-run status. Finds the run whose status is failed.

2

## `get_workflow_run(run_id)`

Full step-by-step trace. Identifies the single failing step.

3

## `get_step_output(run_id, step_id)`

Verbatim error payload from the failing step.

4

## `suggest_retry(run_id)`

Canonical fix plus the retry command.

*n8n-style MCP surface. One env var (RPA\_API\_URL, RPA\_API\_TOKEN) swaps the stub for a real orchestrator: UiPath, n8n, or self-hosted.*

# Verbatim, never paraphrased

## EVIDENCE:

```
{"error": "channel_not_found", "ts": "1747559640.000100", "channel_attempted": "#new-hires"}
```

## REMEDIATION:

replace channel name "#new-hires" with channel ID C09NEW123HIRE, then retry the run

*The end-to-end smoke test re-runs the agent and asserts the exact string **"error": "channel\_not\_found"** appears in EVIDENCE, character for character.*

WHY IT FITS

# Built for the Maestro BPMN track

1

## Inside the platform

Runs end to end in UiPath: a Maestro BPMN Service Task invokes the agent, and the LLM Gateway is on the hot path for every call.

2

## A 1:1 track fit

A BPMN Service Task calling a coded agent is exactly the track brief. Maestro Case is the documented fallback.

3

## Stub-first design

Judges reproduce the full demo with zero orchestrator setup. One env var flips the stub to a live orchestrator.

## WHAT'S NEXT

# Where gemini-rpa-agent goes from here

### Test Cloud surface

A second MCP surface for the third AgentHack track: `list_test_runs`, `get_fragile_tests`, `suggest_test_fix`.

### Human Task approval

A Maestro Human Task fork after diagnosis, so a person approves or rejects the fix before it auto-applies.

### Watch-mode daemon

Poll `list_workflows(active_only)` and only invoke Gemini on a new failed run, which pays for itself in gateway cost.

**Diagnosis in seconds.  
The fix already quoted.  
Every call through the UiPath LLM Gateway.**

[github.com/MukundaKatta/gemini-rpa-agent](https://github.com/MukundaKatta/gemini-rpa-agent)

Maestro BPMN

Gemini 2.5

ADK

MCP

Apache-2.0